

Table 1: Additional scoring alternatives. Each entry uses the same eight pinned targets and six victims per target within an attack (48 trials); $\epsilon = 8/255$. ASR (%) on CIFAR-10 using ConvNetBN, with dog as the poison class and bird as the target class. BP denotes the code’s f c attack in sample mode. Best and second-best values per column are **bold** and underlined, including ties.

Method ρ (%)	ConvNet							
	BP		GM			SAPA		
	0.1	0.2	0.1	0.2	0.5	0.1	0.2	0.5
Classifier-only alignment	<u>29.2</u>	62.5	<u>25.0</u>	<u>29.2</u>	<u>66.7</u>	16.7	31.3	<u>66.7</u>
Normalized classifier alignment	27.1	<u>64.6</u>	29.2	<u>29.2</u>	75.0	<u>22.9</u>	<u>37.5</u>	62.5
Target-class margin	25.0	45.8	20.8	33.3	60.4	25.0	39.6	68.8
Similarity + target-class margin	25.0	68.8	22.9	33.3	54.2	16.7	<u>37.5</u>	52.1
Confidence-only	31.3	47.9	22.9	14.6	54.2	12.5	31.3	64.6

Table 2: GM on a fixed TinyImageNet-20 subset: 20 classes, 200 training images per class (4,000 total), and 1,000 validation images. ResNet18 with a small-image stem is trained from scratch for 30 epochs at 64×64 . Three clean surrogates are shared across methods; $\rho = 0.2\%$ (8 poisons), $\epsilon = 8/255$. Each method uses the same eight targets and six victim seeds per target. ASR is the mean $\pm$ population standard deviation across eight per-target success rates. Accuracy is mean clean validation accuracy after poisoned training. Best and second-best ASR are **bold** and underlined.

Method	ASR (%)	Successes	Accuracy (%)
RAND	2.1 $\pm$ 5.5	1/48	53.64
<i>M</i>	<u>0.0</u> $\pm$ 0.0	0/48	53.50
BASIS	<u>0.0</u> $\pm$ 0.0	0/48	53.30

Table 3: Extended scoring comparison. The first block reproduces the existing ablation; these baselines were not rerun under the new 8×6 protocol. The second block reports the new 48-trial results. Cross-block comparisons are descriptive, not matched reruns. ASR (%) on CIFAR-10 using ConvNetBN, with dog as the poison class and bird as the target class. BP denotes the code’s fc attack in sample mode. Best and second-best values per column are **bold** and underlined, including ties.

Method ρ (%)	ConvNet							
	BP		GM			SAPA		
	0.1	0.2	0.1	0.2	0.5	0.1	0.2	0.5
<i>Previously reported ablation (not rerun)</i>								
RAND	0.0	1.7	0.0	0.0	1.7	0.0	0.0	6.7
R	2.5	22.5	7.5	15.0	20.0	5.0	15.0	25.0
M	<u>35.0</u>	45.0	17.5	20.8	55.0	15.0	27.5	62.5
$\langle g_i, g_t \rangle$	42.5	47.5	15.0	<u>35.0</u>	87.5	20.0	40.0	82.5
BASIS	32.5	70.0	20.0	40.0	<u>77.5</u>	27.5	30.0	62.5
<i>New experiments (8 targets $\times$ 6 victims)</i>								
Classifier-only alignment	29.2	62.5	<u>25.0</u>	29.2	66.7	16.7	31.3	66.7
Normalized classifier alignment	27.1	64.6	29.2	29.2	75.0	22.9	37.5	62.5
Target-class margin	25.0	45.8	20.8	33.3	60.4	<u>25.0</u>	<u>39.6</u>	<u>68.8</u>
Similarity + target-class margin	25.0	<u>68.8</u>	22.9	33.3	54.2	16.7	37.5	52.1
Confidence-only	31.3	47.9	22.9	14.6	54.2	12.5	31.3	64.6